

Endianness

Systems Software

Addresses of multi-byte objects

- Ex: consider the declaration

```
int x;
```

on a 32-bit machine. **x** takes up 4 bytes in memory.


- So if **x** occupies the addresses 0x100 through 0x103, what is the address of **x** (value of the C expression `&x`) ?
- Convention: the address of a multi-byte object is the smallest of the addresses used to store the bytes.

Byte Ordering

- Two popular conventions
- Intel is a little endian architecture
 - Least significant byte is stored at lowest address
- Some other systems use big endian
 - Most significant byte is stored at lowest address
- Ex: How would int x = 0x01234567 at address 0x100 be stored?

Little Endian

- Ex: How would int x = 0x01234567 at address 0x100 be stored?



Origin of Big & Little Endians

Jonathan Swift, *Gulliver's Travels*

- A law requiring all citizens of Lilliput to break their soft-eggs at the little ends only
- A civil war breaking between the Little Endians and the Big Endians, resulting in the Big Endians taking refuge on a nearby island, the kingdom of Blefuscus
- Satire over holy wars between Protestant Church of England and the Catholic Church of France

Do we care?

- For most programs we don't have to worry about endianness ... variables generally work as you would expect


- Endianness matters

- When transferring data over a network between machines with different endianness
- When inspecting binary data representing integers. Consider the following IA32 instruction and its disassembly:

```
80483bd: 01 05 64 94 04 08      add %eax, 0x08049464
```

Pointers

- A pointer is a variable declared to store a memory address
- declaration syntax: `type *var_name;`
- Pointers can be to any data type, including structures



Pointers

■ Important pointer-related operators

`&` : address-of

`*` : dereference (*not the same as the `*` used for declarations!!*)

■ Example

```
int i = 5;      /* i is an int containing 5 */
int *p;         /* p is a pointer to an int */
int j;          /* j is an uninitialized int */

p = &i;         /* store the address of i in p */
j = *p;         /* store the value p points to into j */
```

Pointers and Arrays

- Array name \approx a pointer to the initial (0th) array element
- Example

```
int a[10];
```

```
int *ptr;
```

```
/* following statements are same */
```

```
ptr = a;
```

```
ptr = &a[0];
```

The sizeof operator

- The **sizeof** operator gets a variable or a type as an input and outputs its size in bytes.
- Typical sizes on a 32-bit machine

`sizeof(char) == 1`

`sizeof(short) == 2`

`sizeof(int) == 4`

`sizeof(int *) == 4`

■ Lookup

- signed/unsigned – int, char, short

Pointer arithmetic

- Addition/subtraction operations on pointers work in multiples of the size of the data type being pointed to
- Example

```
char    *cp1;  
int    *ip1;  
double *dp1;
```

```
cp1++;      /* Increments address by 1 */  
ip1++;      /* Increments address by 4 */  
dp1++;      /* Increments address by 8 */
```

Casting

- Reference an object according to a different data type from which it was created.
- You can cast pointers to be pointers to other types. This is powerful.

Code to print byte representation

```
int main(void) {  
    int i=12345, j;  
    unsigned char *p = (unsigned char *) &i;  
  
    for (j=0; j<4; j++)  
        printf("Offset %d: %x\n", j, p[j]);  
}
```

Output on a
little-endian
machine

Offset 0: 39
Offset 1: 30
Offset 2: 0
Offset 3: 0

Decimal: 12345

Binary: 0011 0000 0011 1001

Hex: 3 0 3 9

show_bytes routine P. 37

```
typedef unsigned char *byte_pointer;

void show_bytes(byte_pointer start, int len) {
    int i;
    for (i = 0; i < len; i++)
        printf(" %.2x", start[i]);
    printf("\n");
}
```

Problem P. 40

Consider the following three calls to `show_bytes`

```
int val = 0x12345678;  
byte_pointer valp = (byte_pointer) &val;  
show_bytes(valp, 1); /* A. */  
show_bytes(valp, 2); /* B. */  
show_bytes(valp, 3); /* C. */
```

Indicate the values that would be printed by each call on a little-endian machine and on a big-endian machine.

Problem P. 100

- Write a procedure `is_little_endian` that will return 1 when compiled and run on a little-endian machine, and will return 0 when compiled and run on a big-endian machine. This program should always run on any machine, regardless of its word size.