

The C - Assembly connection

Systems Software

Displacement addressing

■ Used to access a memory address plus a constant displacement

- `movl disp(reg1), reg2`
 - `reg1, reg2` are registers
 - `disp` is a 32-bit constant displacement value
 - Moves the contents of memory at address $(\text{reg1} + \text{disp})$ to `reg2`

■ Example: `movl 8(%ebp), %edx`

- Moves value at memory address `%ebp + 8` to `%edx`

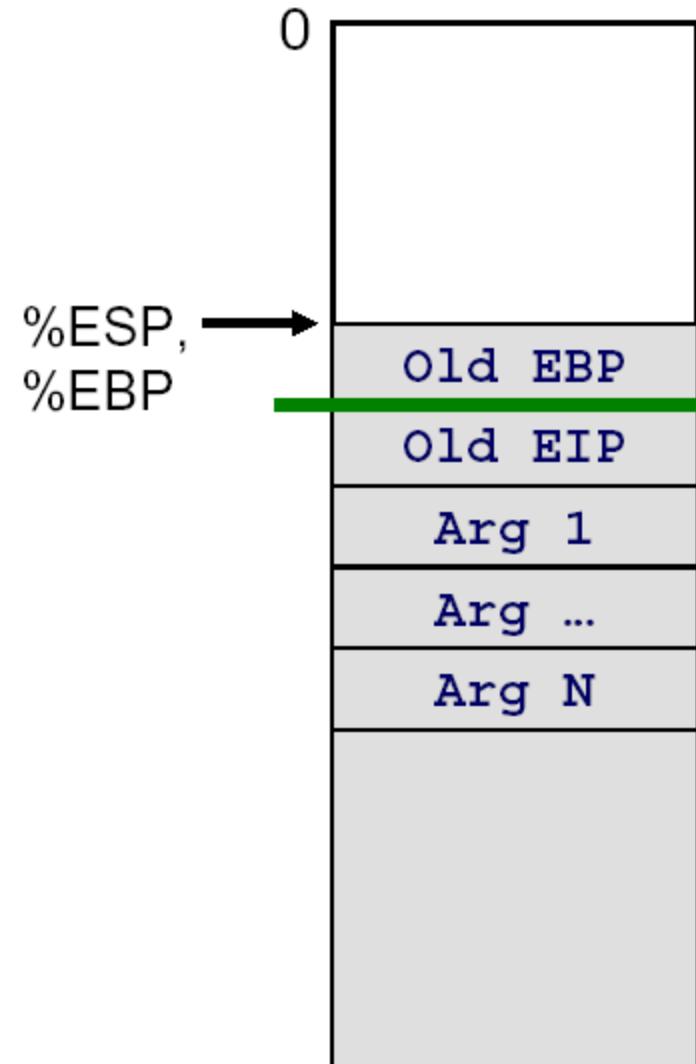
Base Pointer: EBP

- Save caller's base pointer

- `pushl %ebp`

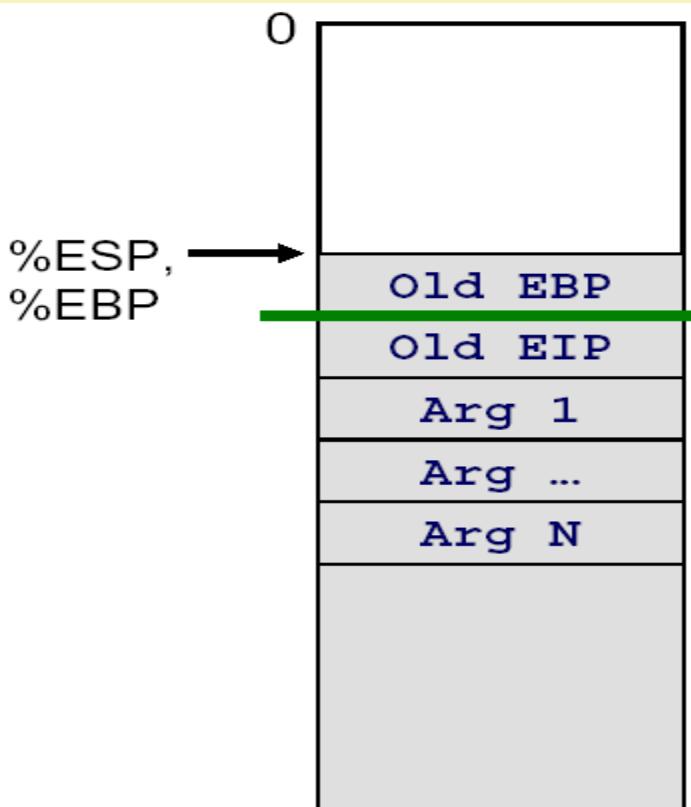
- Set value of EBP for callee's use

- Current value of stack pointer used as callee's base pointer
 - `movl %esp, %ebp`



Using displacement addressing

```
/* sum first 3 elements of  
array a and store result in  
a[3]. */  
  
void sumup(int *a) {  
    int temp;  
    temp = a[0] + a[1] + a[2];  
    a[3] = temp;  
}
```



`sumup:`

```
pushl %ebp  
movl %esp,%ebp  
movl 8(%ebp), %edx
```

} Set Up

```
movl 4(%edx), %eax  
addl (%edx), %eax  
addl 8(%edx), %eax  
movl %eax, 12(%edx)
```

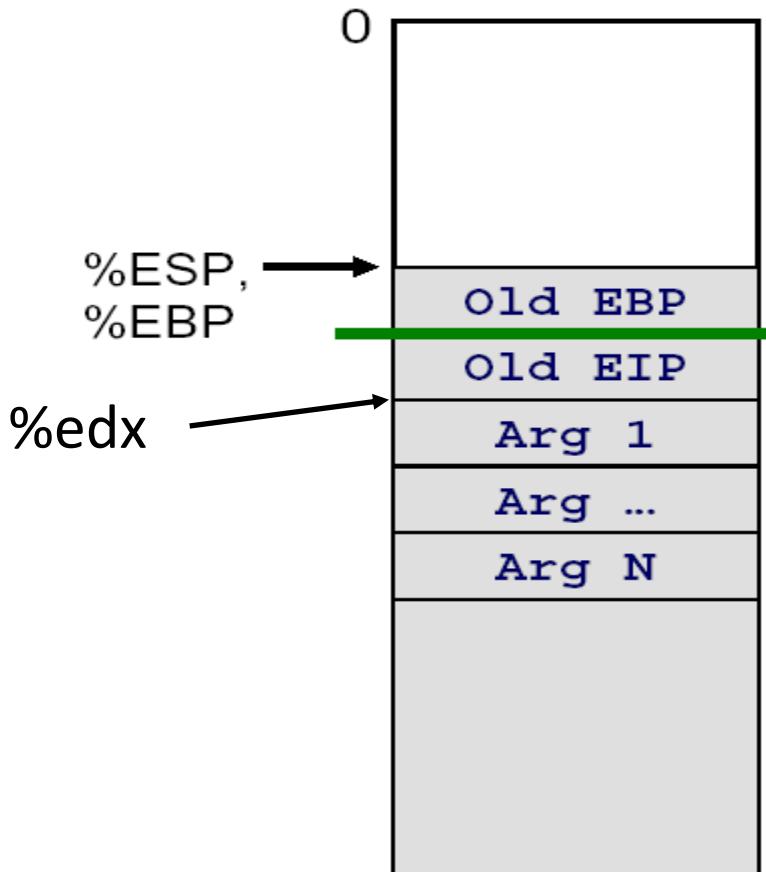
} Body

```
popl %ebp  
ret
```

} Finish

Using displacement addressing

```
void sumup(int *a) {  
    int tmp;  
    temp = a[0] + a[1] + a[2];  
    a[3] = tmp;  
}
```



sumup:

```
pushl %ebp  
movl %esp,%ebp  
movl 8(%ebp), %edx
```

} Set Up

```
movl 4(%edx), %eax  
addl (%edx), %eax  
addl 8(%edx), %eax  
movl %eax, 12(%edx)
```

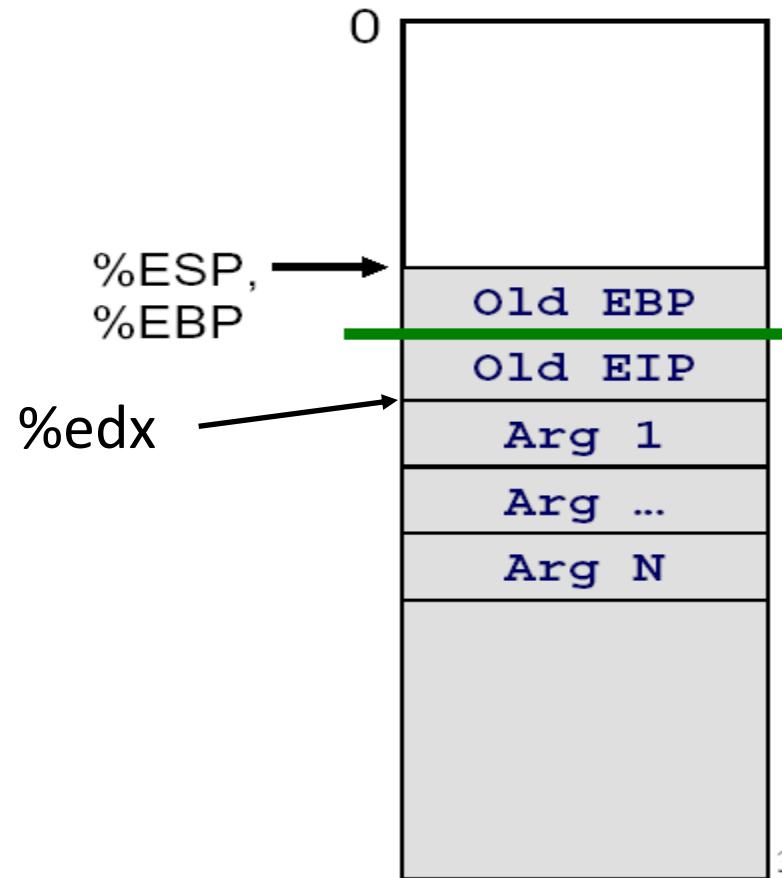
} Body

```
popl %ebp  
ret
```

} Finish

Understanding sumup

```
/* sum first 3 elements
 * of array a and store result
 * in a[3].
 */
void sumup(int *a) {
    int temp;
    temp = a[0] + a[1] + a[2];
    a[3] = temp;
}
```



Value of a (address of first array element) in %edx

```
movl 4(%edx), %eax      ; %eax = a[1]
addl (%edx), %eax       ; %eax += a[0]
addl 8(%edx), %eax      ; %eax += a[2]
movl %eax, 12(%edx)     ; a[3] = %eax
```

Indexed addressing

■ A very general purpose addressing mode

- **disp(reg1, reg2, scale)** accesses memory address: $\text{reg1} + (\text{reg2} * \text{scale}) + \text{disp}$
 - disp: constant displacement (32 bits)
 - reg1: Base register: Any of the 8 general-purpose registers
 - reg2: Index register: Any register except for %esp
 - scale: Must be 1, 2, 4, or 8

■ Example: **movl 10(%eax, %ebx, 4), %ecx**

- Accesses memory at address $\%eax + (\%ebx * 4) + 10$
- Stores result in %ecx

Address Computation Examples

%edx	0xf000
%ecx	0x100

Expression	Computation	Address
0x8(%edx)		
(%edx, %ecx)		
(%edx, %ecx, 4)		
0x80(, %edx, 2)		

Address Computation Examples

%edx	0xf000
%ecx	0x100

Expression	Computation	Address
0x8(%edx)	0xf000 + 0x8	0xf008
(%edx, %ecx)		
(%edx, %ecx, 4)		
0x80(,%edx,2)		

Address Computation Examples

%edx	0xf000
%ecx	0x100

Expression	Computation	Address
0x8(%edx)	0xf000 + 0x8	0xf008
(%edx, %ecx)	0xf000 + 0x100	0xf100
(%edx, %ecx, 4)		
0x80(,%edx,2)		

Address Computation Examples

%edx	0xf000
%ecx	0x100

Expression	Computation	Address
0x8(%edx)	0xf000 + 0x8	0xf008
(%edx, %ecx)	0xf000 + 0x100	0xf100
(%edx, %ecx, 4)	0xf000 + 4*0x100	0xf400
0x80(,%edx,2)		

Address Computation Examples

%edx	0xf000
%ecx	0x100

Expression	Computation	Address
0x8(%edx)	0xf000 + 0x8	0xf008
(%edx, %ecx)	0xf000 + 0x100	0xf100
(%edx, %ecx, 4)	0xf000 + 4*0x100	0xf400
0x80(,%edx,2)	2*0xf000 + 0x80	0x1e080

Computing addresses: `leal` instruction

- Sometimes want to save the result of an address computation for later use
 - Rather than reading or writing memory at the computed address
- The `leal` instruction is used for this: “load effective address”
- Example: `leal 4(%eax, %ebx, 8), %ecx`
 - Computes the address of the “source” argument, stores in the destination
 - Does not access the memory address itself: use `movl` for that!
- Can also be used to compute arithmetic expressions!
 - Anything of the form $x + y*k$, where $k = 1, 2, 4$, or 8

A more complex example

■ What does this code do?

```
# Address of array in %ecx
movl 8(%ecx), %eax
movl 4(%ecx), %edx
leal (%edx, %eax, 2), %eax
movl (%ecx), %edx
leal (%edx, %eax, 2), %eax
movl %eax, 12(%ecx)
```

A more complex example

■ What does this code do?

```
# Address of array[] in %ecx
movl 8(%ecx), %eax          ; eax = array[2]
movl 4(%ecx), %edx          ; edx = array[1]
leal (%edx, %eax, 2), %eax ; eax = array[1] + (2*array[2])
movl (%ecx), %edx           ; edx = array[0]
leal (%edx, %eax, 2), %eax ; eax = array[0] + (2 * eax)
movl %eax, 12(%ecx)         ; array[3] = eax
```

■ Original source code

```
void dosomething(int *array) {
    int tmp;
    tmp = array[0] + (2 * array[1]) + (4 * array[2]);
    array[3] = tmp;
}
```

Structures

```
struct rec {  
    int i;  
    int a[3];  
    int *p;  
};
```

Structures

```
struct rec {  
    int i;  
    int a[3];  
    int *p;  
};
```

Memory Layout



Characteristics

- Contiguously-allocated region of memory
- Refer to members within structure by names
- Members may be of different types

Structures

■ Accessing Structure Member

- Given an instance of the struct, we can use the `.` operator:
 - `struct rec r1; r1.i = val;`
- What if we have a *pointer* to a struct: `struct rec *r = &r1;`

```
struct rec {  
    int i;  
    int a[3];  
    int *p;  
};
```

Structures

■ Accessing Structure Member

- Given an instance of the struct, we can use the `.` operator, just like Java:
 - `struct rec r1; r1.i = val;`
- What if we have a *pointer* to a struct: `struct rec *r = &r1;`
 - Using `*` and `.` operators: `(*r).i = val;`
 - Or, use `->` operator for short: `r->i = val;`
- Pointer indicates first byte of structure; access members with offsets

```
void
set_i(struct rec *r,
      int val)
{
    r->i = val;
}
```

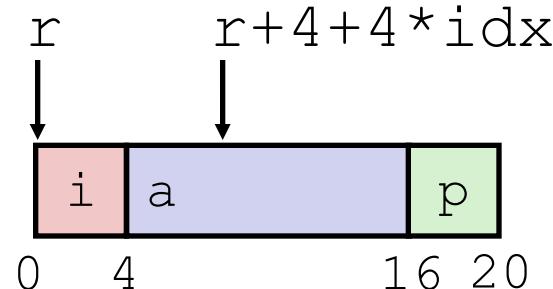
```
struct rec {
    int i;
    int a[3];
    int *p;
};
```

IA32 Assembly

```
# %eax = val
# %edx = r
movl %eax, (%edx)    # Mem[r] = val
```

Generating Pointer to Structure Member

```
struct rec {  
    int i;  
    int a[3];  
    int *p;  
};
```



■ Generating Pointer to Array Element

- Offset of each structure member determined at compile time

```
int *find_a  
(struct rec *r, int idx)  
{  
    return &r->a[idx];  
}
```

```
# %ecx = idx  
# %edx = r  
leal 0(,%ecx,4),%eax    # 4*idx  
leal 4(%eax,%edx),%eax # r+4*idx+4
```

Practice question

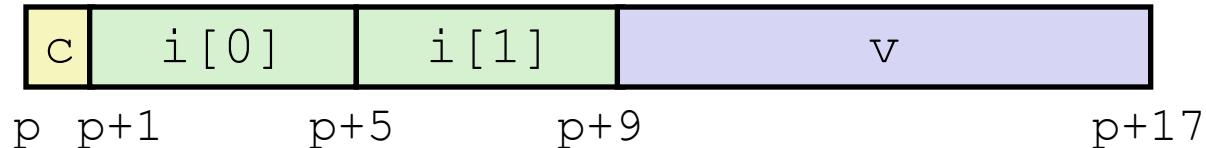
```
struct prob {  
    int *p;  
    struct {  
        int x;  
        int y;  
    } s;  
    struct prob *next;  
};
```

```
void sp_init(struct prob *sp)  
{  
    sp->s.x = _____;  
    sp->p = _____;  
    sp->next = _____;  
}
```

```
1      movl 8(%ebp), %eax  
2      movl 8(%eax), %edx  
3      movl %edx, 4(%eax)  
4      leal 4(%eax), %edx  
5      movl %edx, (%eax)  
6      movl %eax, 12(%eax)
```

Structures & Alignment

■ Unaligned Data

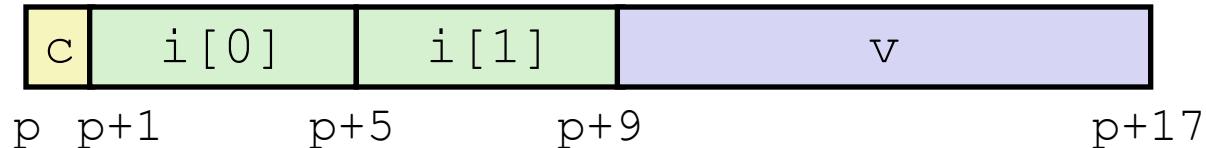


```
struct S1 {  
    char c;  
    int i[2];  
    double v;  
} *p;
```

- How would it look like if data items were *aligned (address multiple of type size)*?

Structures & Alignment

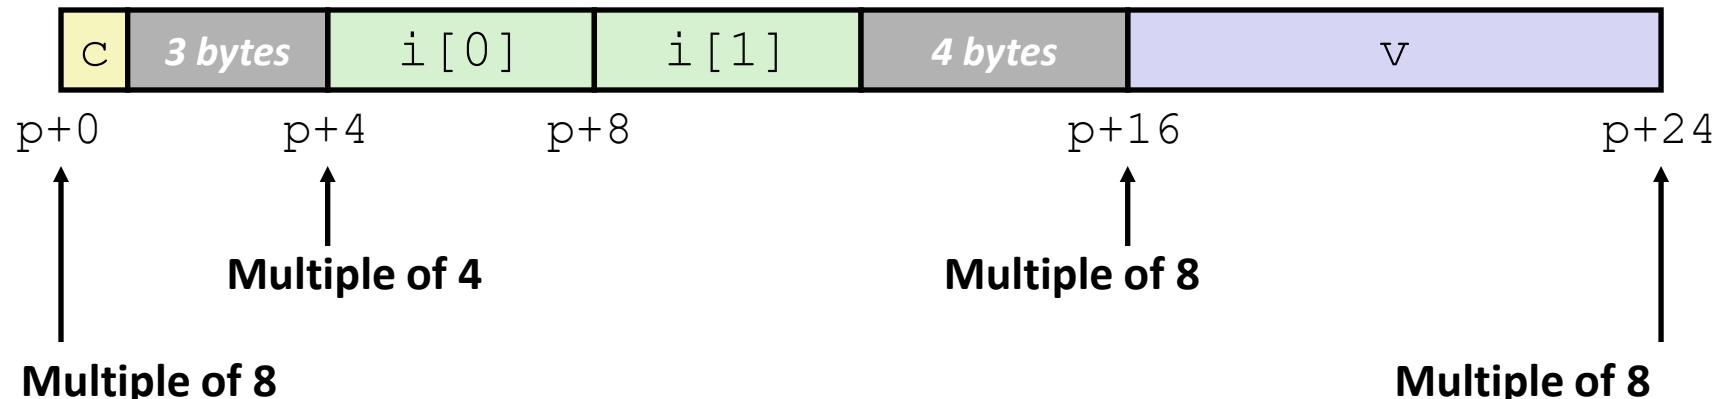
■ Unaligned Data



```
struct S1 {  
    char c;  
    int i[2];  
    double v;  
} *p;
```

■ Aligned Data

- Primitive data type requires K bytes
- Address must be multiple of K



Alignment Principles

■ Aligned Data

- Primitive data type requires K bytes
- Address must be multiple of K

■ Aligned data is required on some machines; it is *advised* on IA32

- Treated differently by IA32 Linux, x86-64 Linux, and Windows!

■ What is the motivation for alignment?

Alignment Principles

■ Motivation for Aligning Data

- Physical memory is accessed by aligned chunks of 4 or 8 bytes (system-dependent)
 - Inefficient to load or store datum that crosses quad word boundaries

■ Compiler

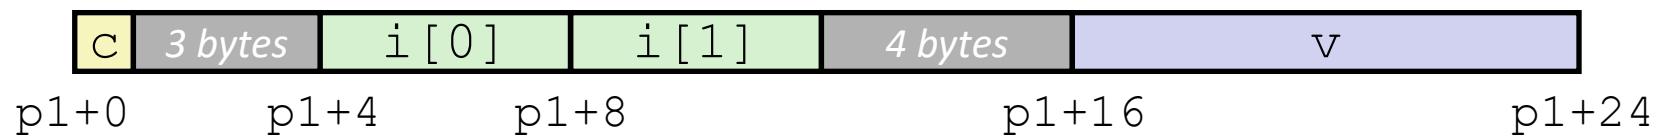
- Inserts padding in structure to ensure correct alignment of fields
- `sizeof()` should be used to get true size of structs

Specific Cases of Alignment (IA32)

- **1 byte: char, ...**
 - no restrictions on address
- **2 bytes: short, ...**
 - lowest 1 bit of address must be 0_2
- **4 bytes: int, float, char *, ...**
 - lowest 2 bits of address must be 00_2
- **8 bytes: double, ...**
 - Windows (and most other OSs & instruction sets): lowest 3 bits 000_2
 - Linux: lowest 2 bits of address must be 00_2
 - i.e., treated the same as a 4-byte primitive data type

Saving Space

```
struct S1 {  
    char c;  
    int i[2];  
    double v;  
} *p1;
```



Saving Space

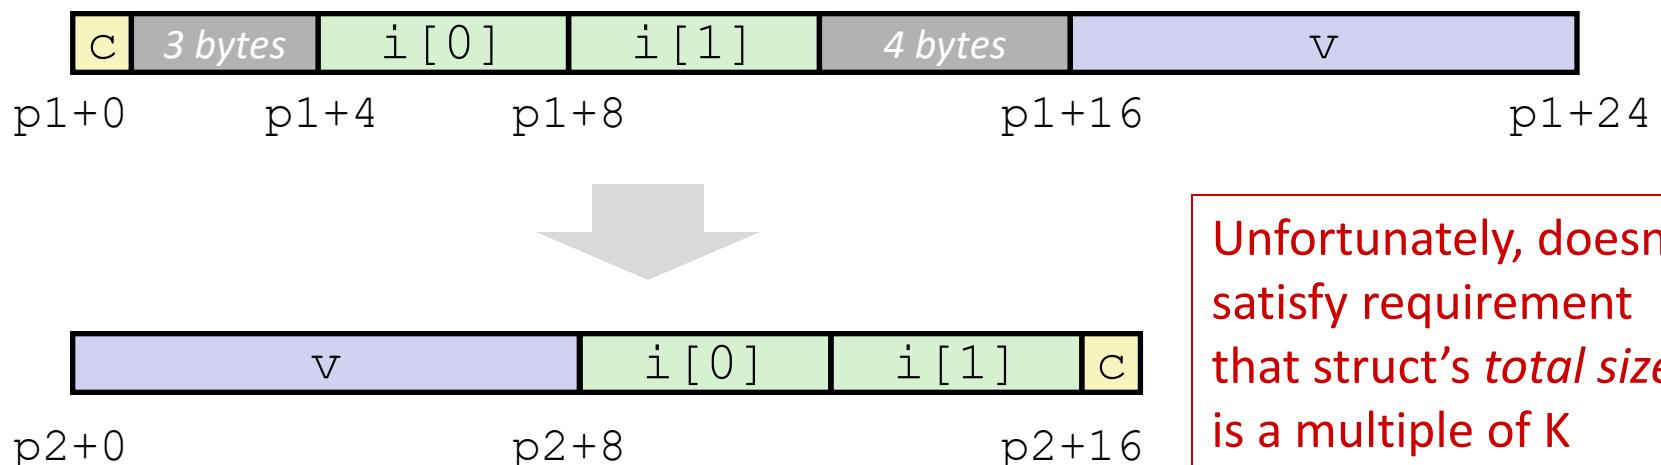
■ Put large data types first:

```
struct S1 {  
    char c;  
    int i[2];  
    double v;  
} *p1;
```



```
struct S2 {  
    double v;  
    int i[2];  
    char c;  
} *p2;
```

■ Effect (both have K=8)



Unfortunately, doesn't satisfy requirement that struct's *total size* is a multiple of K

Arrays of Structures

- Satisfy alignment requirement for every element
- How would accessing an element work?

```
struct S2 {  
    double v;  
    int i[2];  
    char c;  
} a[10];
```

