

The C - Assembly connection

Systems Software

Intel IA32 Architecture

- Often referred to as “x86”

- Started with Intel 80386 in 1985 ... then 80486, Pentium, Intel Core, Intel Core 2, ---
- The most popular ISA by far

- Has a large and varied set of instructions

- Fortunately we won't need to use everything!

Compiling into assembly

■ C Code

```
int sum(int x, int y)
{
    int t = x+y;
    return t;
}
```

Generated IA32 Assembly

```
_sum:
    pushl %ebp
    movl %esp,%ebp
    movl 12(%ebp),%eax
    addl 8(%ebp),%eax
    movl %ebp,%esp
    popl %ebp
    ret
```

Obtain with command

```
gcc -O -S code.c
```

Produces file code.s

Disassembling Object Code

```
00401040 <_sum>:  
 0:      55          push    %ebp  
 1:  89 e5        mov     %esp, %ebp  
 3:  8b 45 0c    mov     0xc(%ebp), %eax  
 6:  03 45 08    add     0x8(%ebp), %eax  
 9:  89 ec        mov     %ebp, %esp  
 b:  5d          pop    %ebp  
 c:  c3          ret
```

■ Disassembler: `objdump -d code.o`

- Produces interpretation of assembly code
- Can be run on either `a.out` (complete executable) or `.o` file

Alternate Disassembly

Object

```
0x401040: 0x55  
           0x89  
           0xe5  
           0x8b  
           0x45  
           0x0c  
           0x03  
           0x45  
           0x08  
           0x89  
           0xec  
           0x5d  
           0xc3
```

Disassembled

0x401040 <sum>:	push	%ebp
0x401041 <sum+1>:	mov	%esp, %ebp
0x401043 <sum+3>:	mov	0xc(%ebp), %eax
0x401046 <sum+6>:	add	0x8(%ebp), %eax
0x401049 <sum+9>:	mov	%ebp, %esp
0x40104b <sum+11>:	pop	%ebp
0x40104c <sum+12>:	ret	

■ Within gdb Debugger

`gdb code`

`disassemble sum`

- Disassemble procedure

`x/13b sum`

- Examine the 13 bytes starting at `sum`

Some x86 registers

31	15	8	7	0	16-bit	32-bit	64-bit
		AH	AL		AX	EAX	RAX
		BH	BL		BX	EBX	RBX
		CH	CL		CX	ECX	RCX
		DH	DL		DX	EDX	RDX
			SI			ESI	
			DI			EDI	
							...

General-purpose registers

GNU Assembly Syntax

GDB: GNU Debugger

GCC: GNU Compiler Collection

GAS: GNU Assembler

source first

Meaning	GAS
ebx := eax	movl %eax, %ebx
eax := eax + ebx	addl %ebx, %eax
ecx := ecx << 2	shl \$2, %ecx

- Referring to a register: percent sign (“%”)
 - E.g., “%ecx”
- Referring to a constant: dollar sign (“\$”)
 - E.g., “\$1” for the number 1

Addressing modes

- Most instructions have several ways of addressing source and destination operands
 - Inputs can be registers, memory location, or immediate (constant) values
 - Outputs can be saved to registers or memory locations
 - Example: “**movl**” instruction (copy 32-bit values) supports...
 - Immediate to register `movl $0x1000, %eax`
 - Register to Register `movl %eax, %ebx`
 - Memory to register (a.k.a. “load”) `movl (%eax), %ebx`
 - Register to memory (a.k.a. “store”) `movl %eax, (%ebx)`

Memory references

Addresses are indicated by operands that have a paren “()”

Register	Value
eax	0x3
edx	0x0
ebx	0x5

What does
mov (%al), %dl
do?

Moves 0xcc
into dl

0xff	6
0xee	
0xdd	
0xcc	
0xbb	
0xaa	
0x00	0

Memory references

Addresses are indicated by operands that have a paren “()”

Register	Value
eax	0x3
edx	0xcc
ebx	0x5

What does
mov (%eax), %edx
do?

Which 4 bytes get moved, and which is the LSB in edx?

0xff	6
0xee	
0xdd	
0xcc	
0xbb	
0xaa	
0x00	0

`mov (%eax), %edx`

EDX

Register	Value
<code>eax</code>	<code>0x3</code>
<code>edx</code>	<code>0xcc</code>
<code>ebx</code>	<code>0x5</code>

EDX =
`0xffeeddcc!`

Bit 0

0xff	6
0xee	
0xdd	
0xcc	
0xbb	
0xaa	
0x00	0

Little Endian: Least significant byte first

... so ...

address a goes in the least significant byte
(the **littlest** bit) $a+1$ goes into the next byte,
and so on.

mov %ebx, (%eax)

Register	Value	EBX	00	00	00	05	Bit 0	Addr 6
eax	0x3							
edx	0xcc							
ebx	0x5							
							0xff	
							0xee	
							0xdd	
							0xcc	
							0xbb	
							0xaa	
							0x00	0

Little Endian: Least significant byte first

... so ...

address a goes in the least significant byte
(the **littlest** bit) $a+1$ goes into the next byte,
and so on.

Program Memory Layout

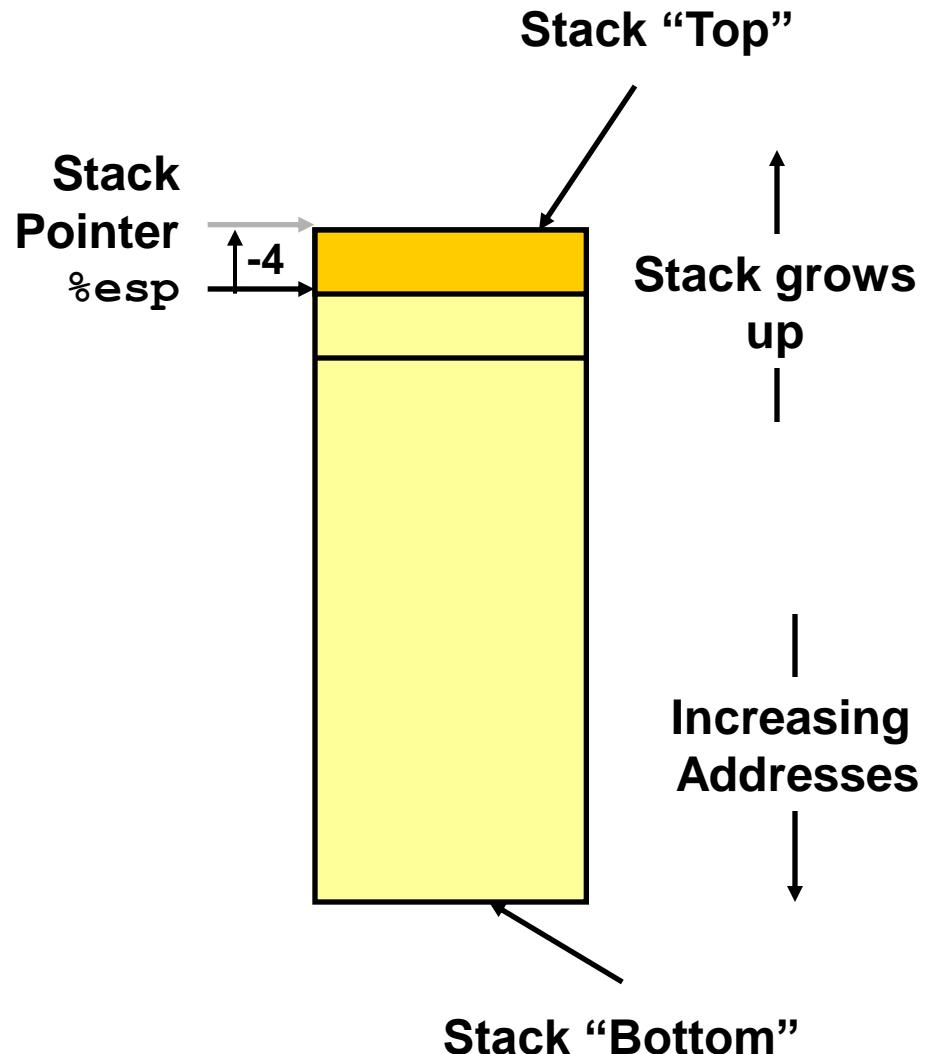
- Text region
 - Program's executable code
- Heap
 - Dynamically allocated data
- Stack
 - Stores all temporary data related to each function call
 - Return address
 - Arguments
 - Local variables of function

Registers for Executing the Code

- Instruction pointer (EIP)
 - Address in memory of the next instruction
- Interesting pointers to the stack
 - Stack register (ESP)
 - Address of the top of the stack
 - Base pointer (EBP)
 - Used for relative references to local variables and arguments

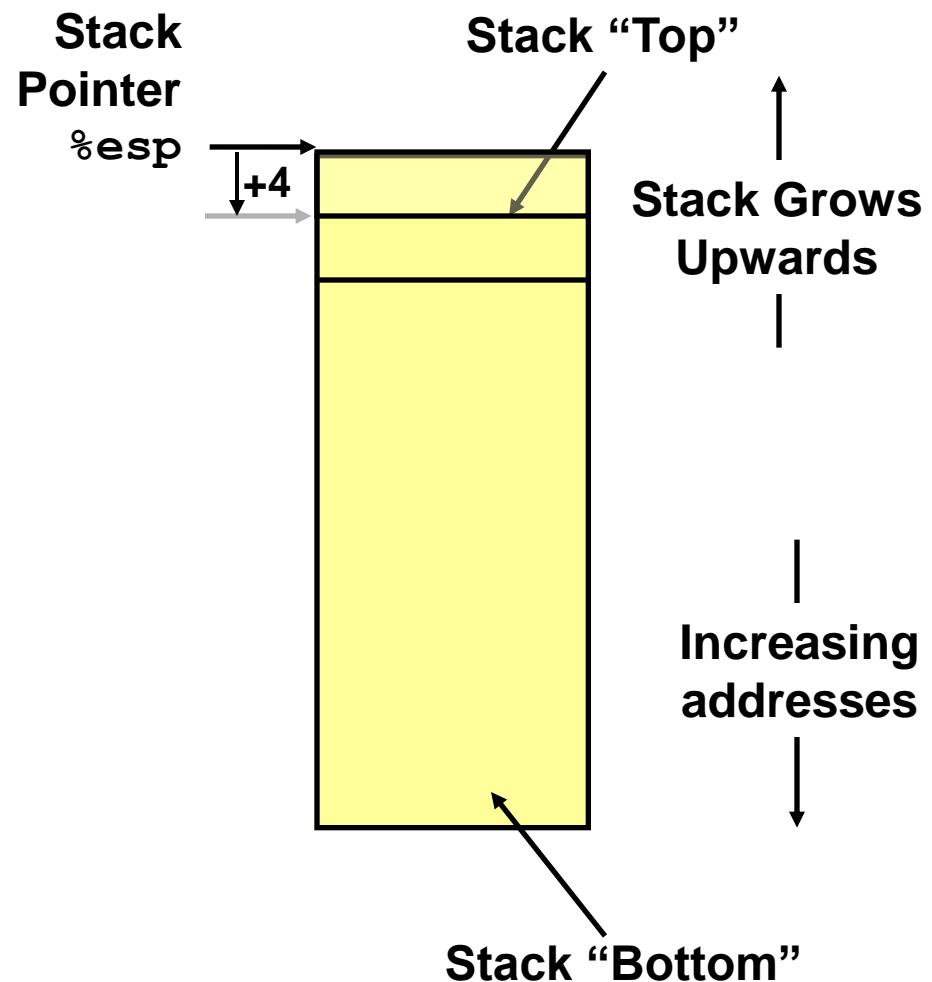
Push Operation

- Push instruction
 - `pushl Src`
- Fetch operand at *Src*
- **Decrement %esp** by 4
- Write operand at address given by %esp



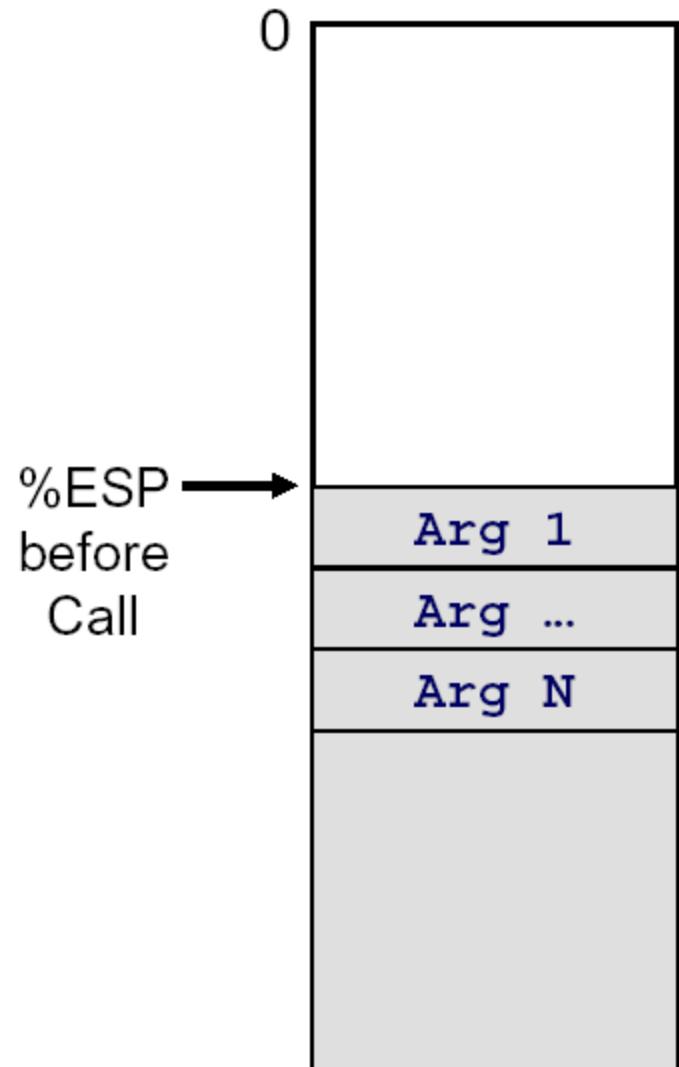
Pop Operation

- Popping
 - `popl Dest`
 - Read operand at address given by `%esp`
 - **Increment** `%esp` by 4
 - Write to *Dest*



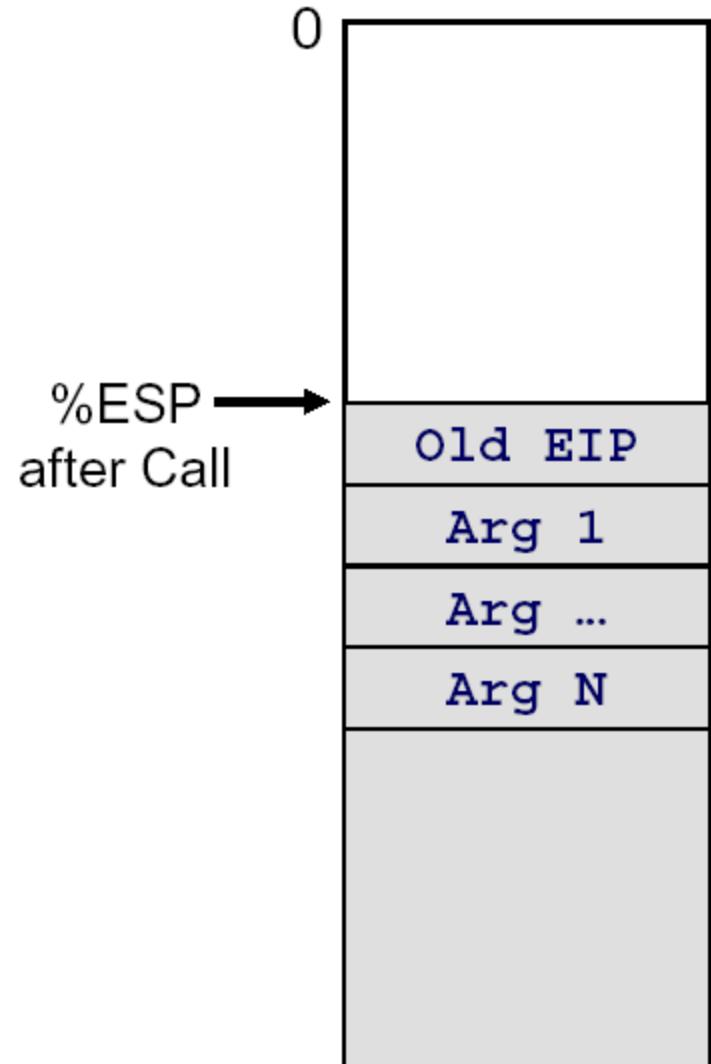
Input parameters

- Caller pushes input parameters before executing call instruction
- Parameters are pushed in reverse order
 - So that the first argument is at the top of the stack at the time of the call
- When done, Caller invokes the function with the “call”



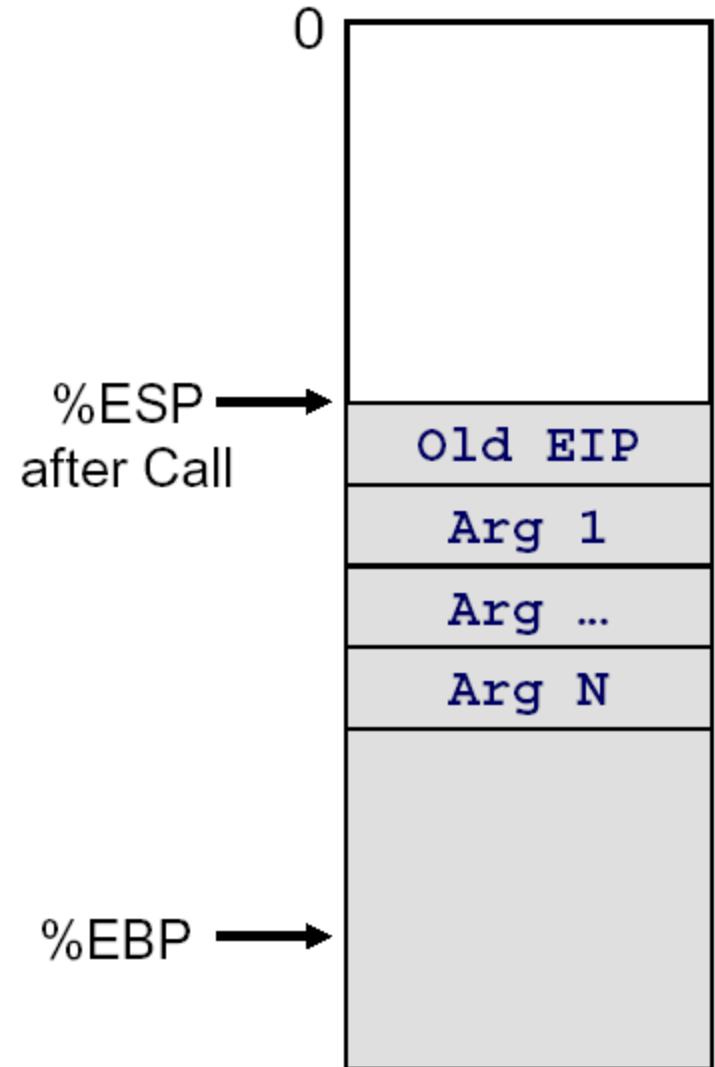
Input parameters

- Call instruction pushes %eip onto stack (return address)
- Call instruction causes execution to shift to the callee
- Callee can address arguments relative to ESP: Arg1 as 4(%esp)



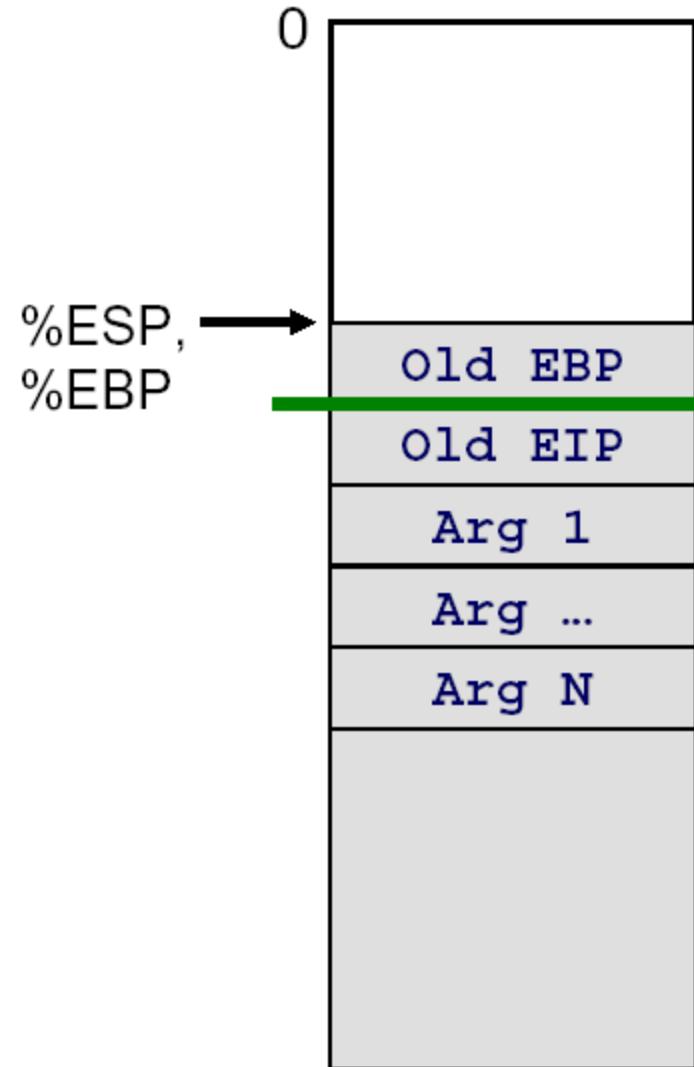
Base Pointer: EBP

- As callee executes, ESP may change
- Use EBP as a fixed reference point to access arguments and other local variables
- Need to save caller's value of EBP before using EBP



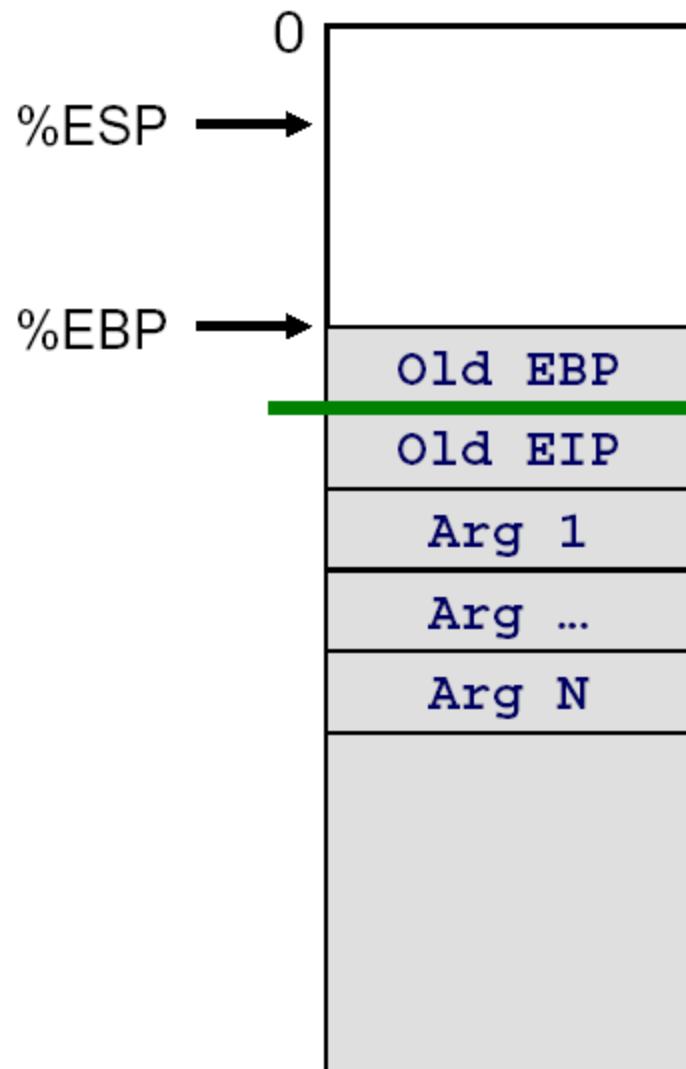
Base Pointer: EBP

- Save caller's base pointer
 - `pushl %ebp`
- Set value of EBP for callee's use
 - Current value of stack pointer used as callee's base pointer
 - `movl %esp, %ebp`



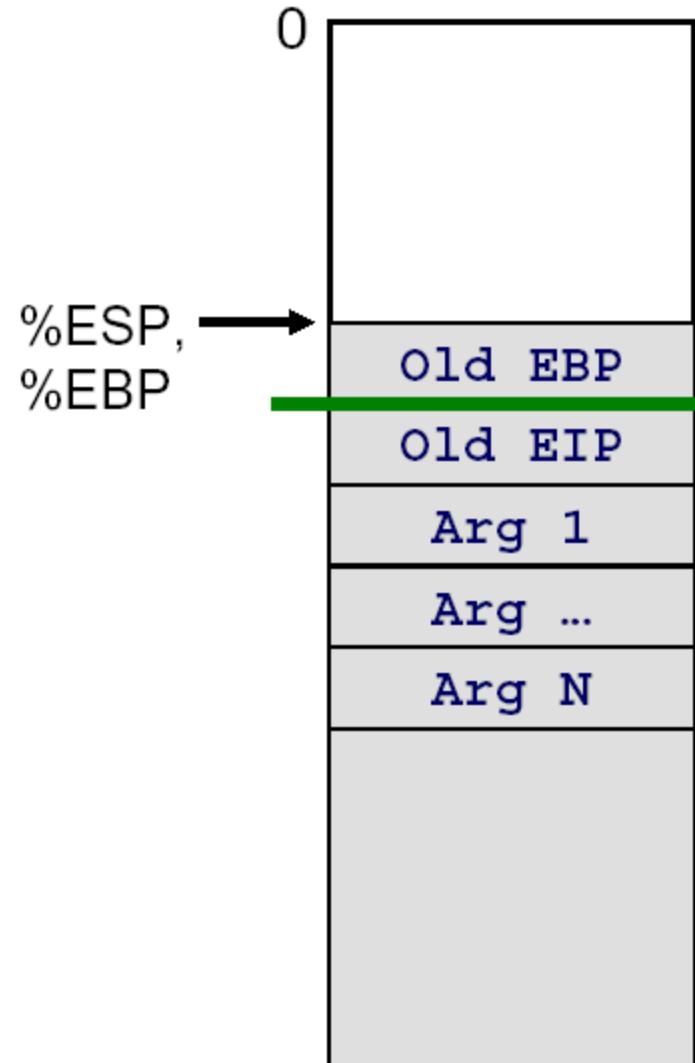
Base Pointer: EBP

- As Callee executes, ESP may change
- Regardless of ESP, Callee can address Arg1 as 8(%ebp)



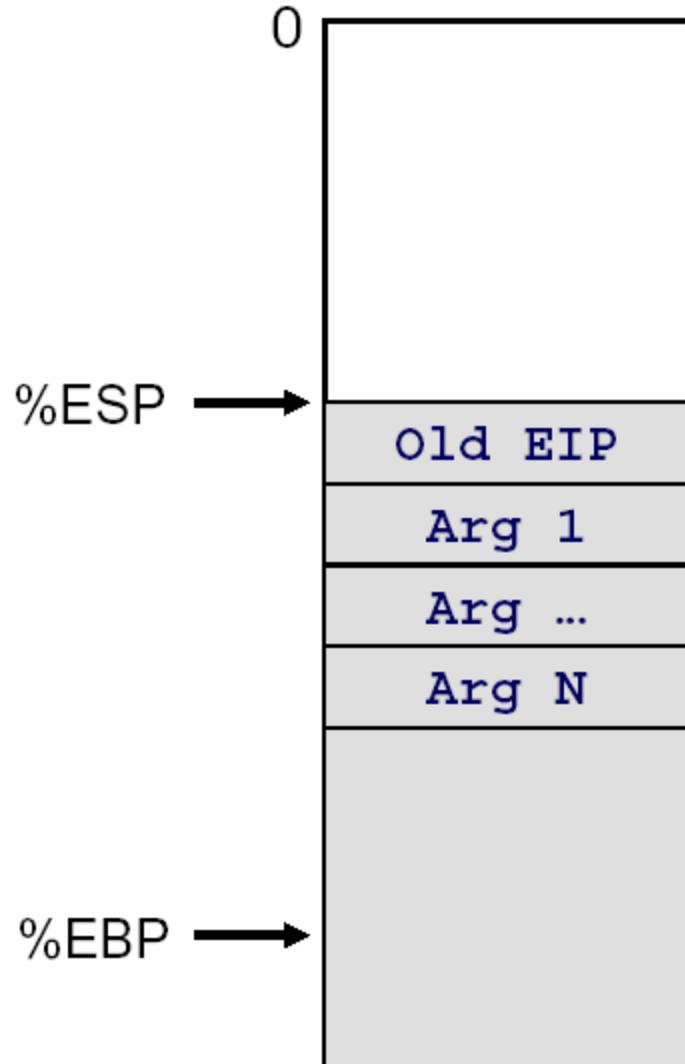
Base Pointer: EBP

- Before returning,
Callee must restore
EBP to its old value
- Executes
 - movl \%ebp, \%esp
 - popl \%ebp
 - ret



Base Pointer: EBP

- Before returning,
Callee must restore
EBP to its old value
- Executes
 - `movl %ebp, %esp`
 - `popl %ebp`
 - `ret`



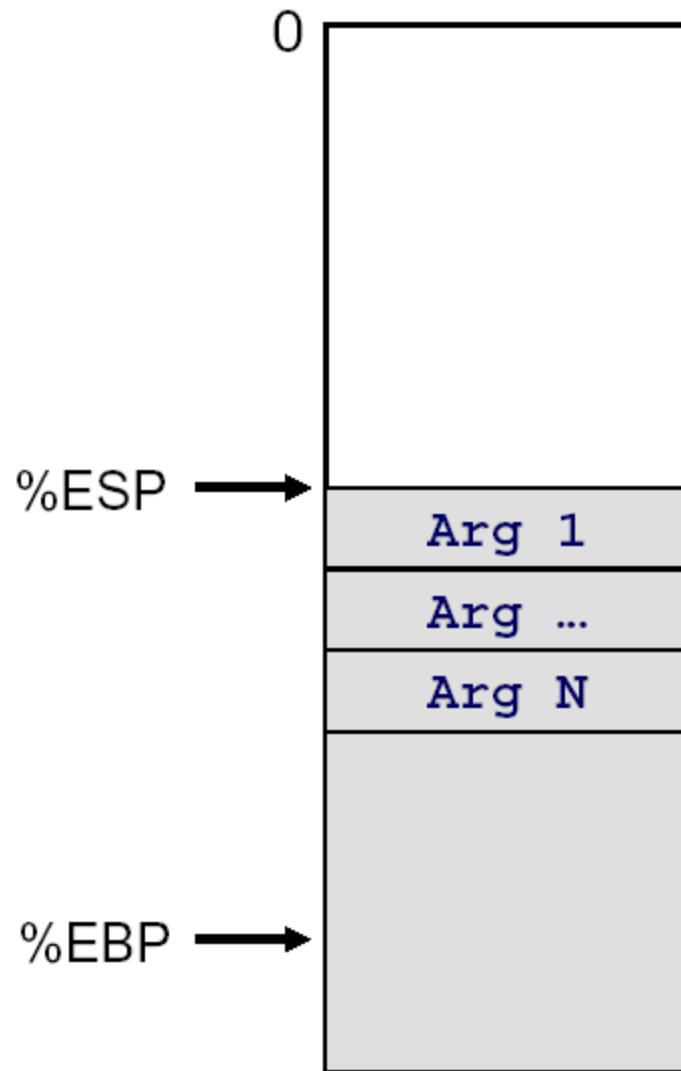
Base Pointer: EBP

- Before returning,
Callee must restore
EBP to its old value
- Executes

```
    movl %ebp, %esp
```

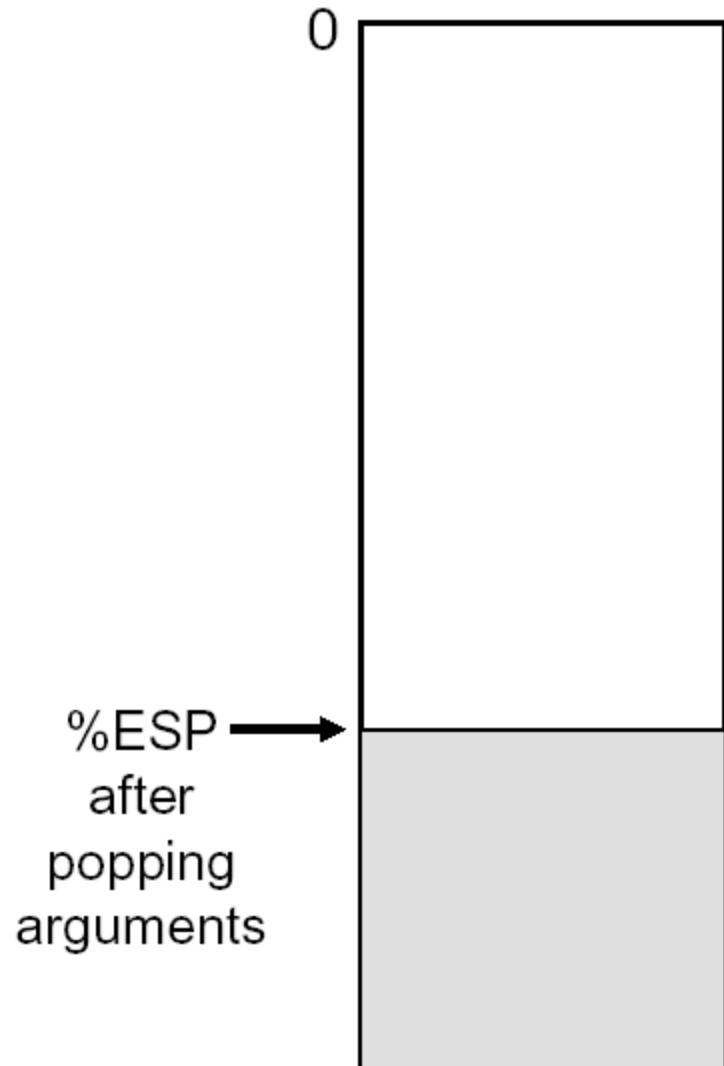
```
    popl %ebp
```

```
    ret
```



Input parameters

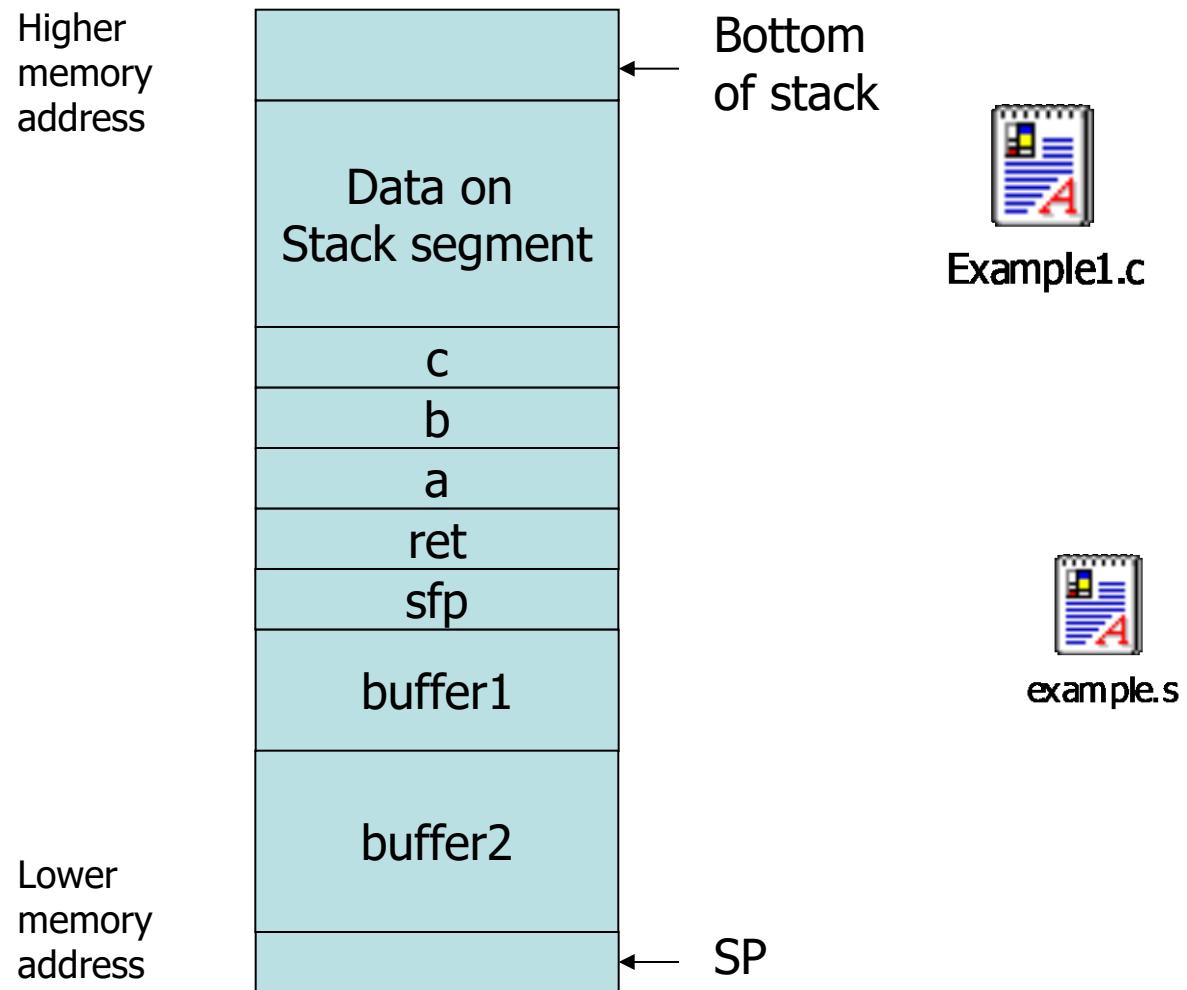
- After the function call is finished, the caller pops the pushed arguments from the stack



Key Point

- Function arguments have fixed positive offsets relative to frame pointer
 - +8(%ebp)
 - +12(%ebp)
- Local variable have fixed negative offsets relative to frame pointer
 - -4(%ebp)
 - -8(%ebp)

Example



Teaser Problem

Practice problem 3.14 from Bryant

You will see this code all over if you look at library routines

```
call next  
next:  
    popl %eax
```

- What value does `%eax` end up with?
- Why is there no `ret` to match the call?
- What useful purpose might this code serve?

Summary

- Invoking a function
 - Call: call the function
 - Ret: return from the instruction
- Stack Frame for a function invocation includes
 - Return address
 - Procedure arguments
 - Local variables
- Base pointer EBP
 - Fixed reference point in the Stack Frame
 - Used for referencing arguments and local variables